

Operating System

lina liu Apr. 13

- 1.how ls cmd process executed?
- 2.what is process?
- 3.what is virtual memory?
- 4.what the difference between memory cache and virtual memory?
- 5.what is CPU protect mode?
- 6.difference between BIOS and UEFI under CPU protect mode?

3 pieces operating system

Virtualization

The Abstraction: The process

Interlude: Process API

Mechanism: Limited Direct Execution

Scheduling: Introduction

Scheduling: The multi-level feedback queue

Scheduling: Introduction

Mechanism: Limited Direct Execution

Scheduling: Proportional Share

Multiprocessor Scheduling

CPU Virtualization

Memory Virtualization

The abstraction: Address Space

Interlude: Memory API

Mechanism: Address Translation

Segmentation

Free-Space Management

Paging: Introduction

Paging: Faster Translation(TLBs)/Smaller tables

Beyond Physical Memory: Mechanisms/Policies

Complete Virtual Memory Systems

Memory Virtualization Summary

Concurrency

Concurrency: An Introduction

Interlude: Thread API

Locks

Lock-based concurrent data structure

Condition Variables

Semaphores

Common Concurrency Problems

Event-based concurrency(Advanced)

Summary of Concurrency

Persistence

A Dialogue on Persistence

I/O Devices

Hard Disk Drives

Redundant Arrays of Inexpensive Disks(RAIDs)

Interlude: Files and Directories

File System Implementation

Locality and the fast file system

Crash Consistency: FSCK and Journaling

Log-structured file systems

Flash-based SSDs

Data Integrity and Protection

Summary Dialogue on Persistence

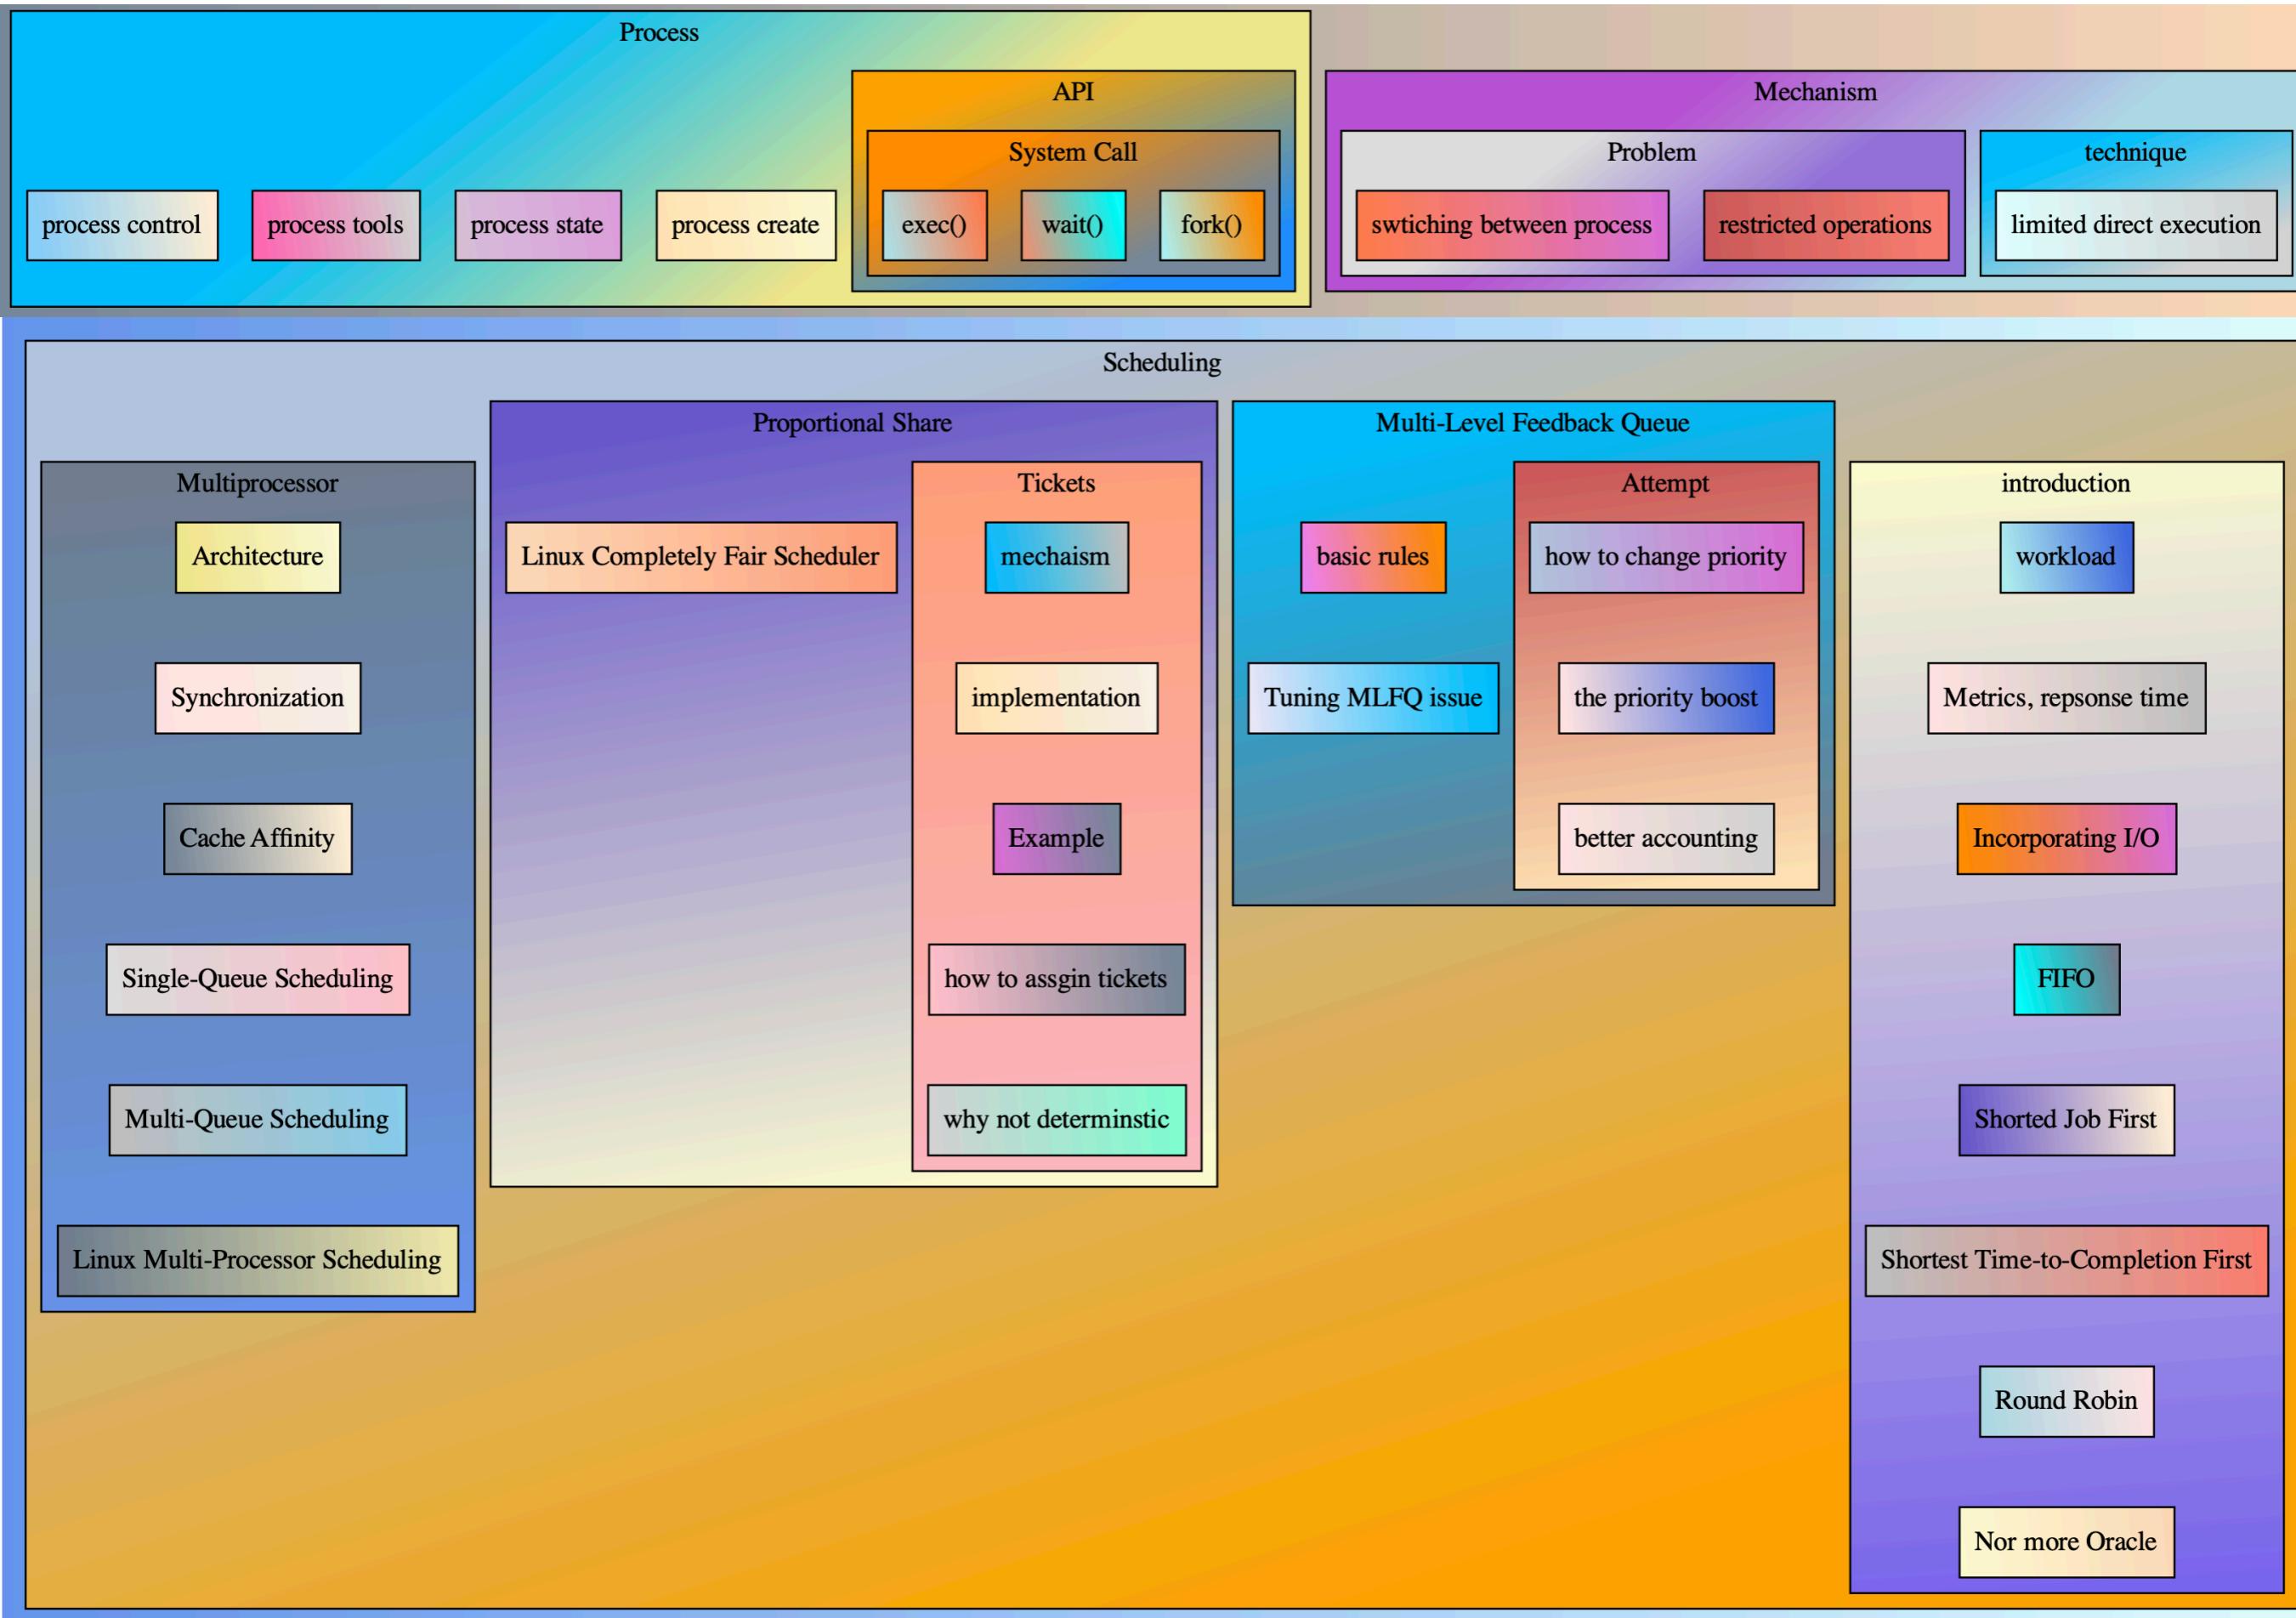
A Dialogue on Distribution

Distributed Systems

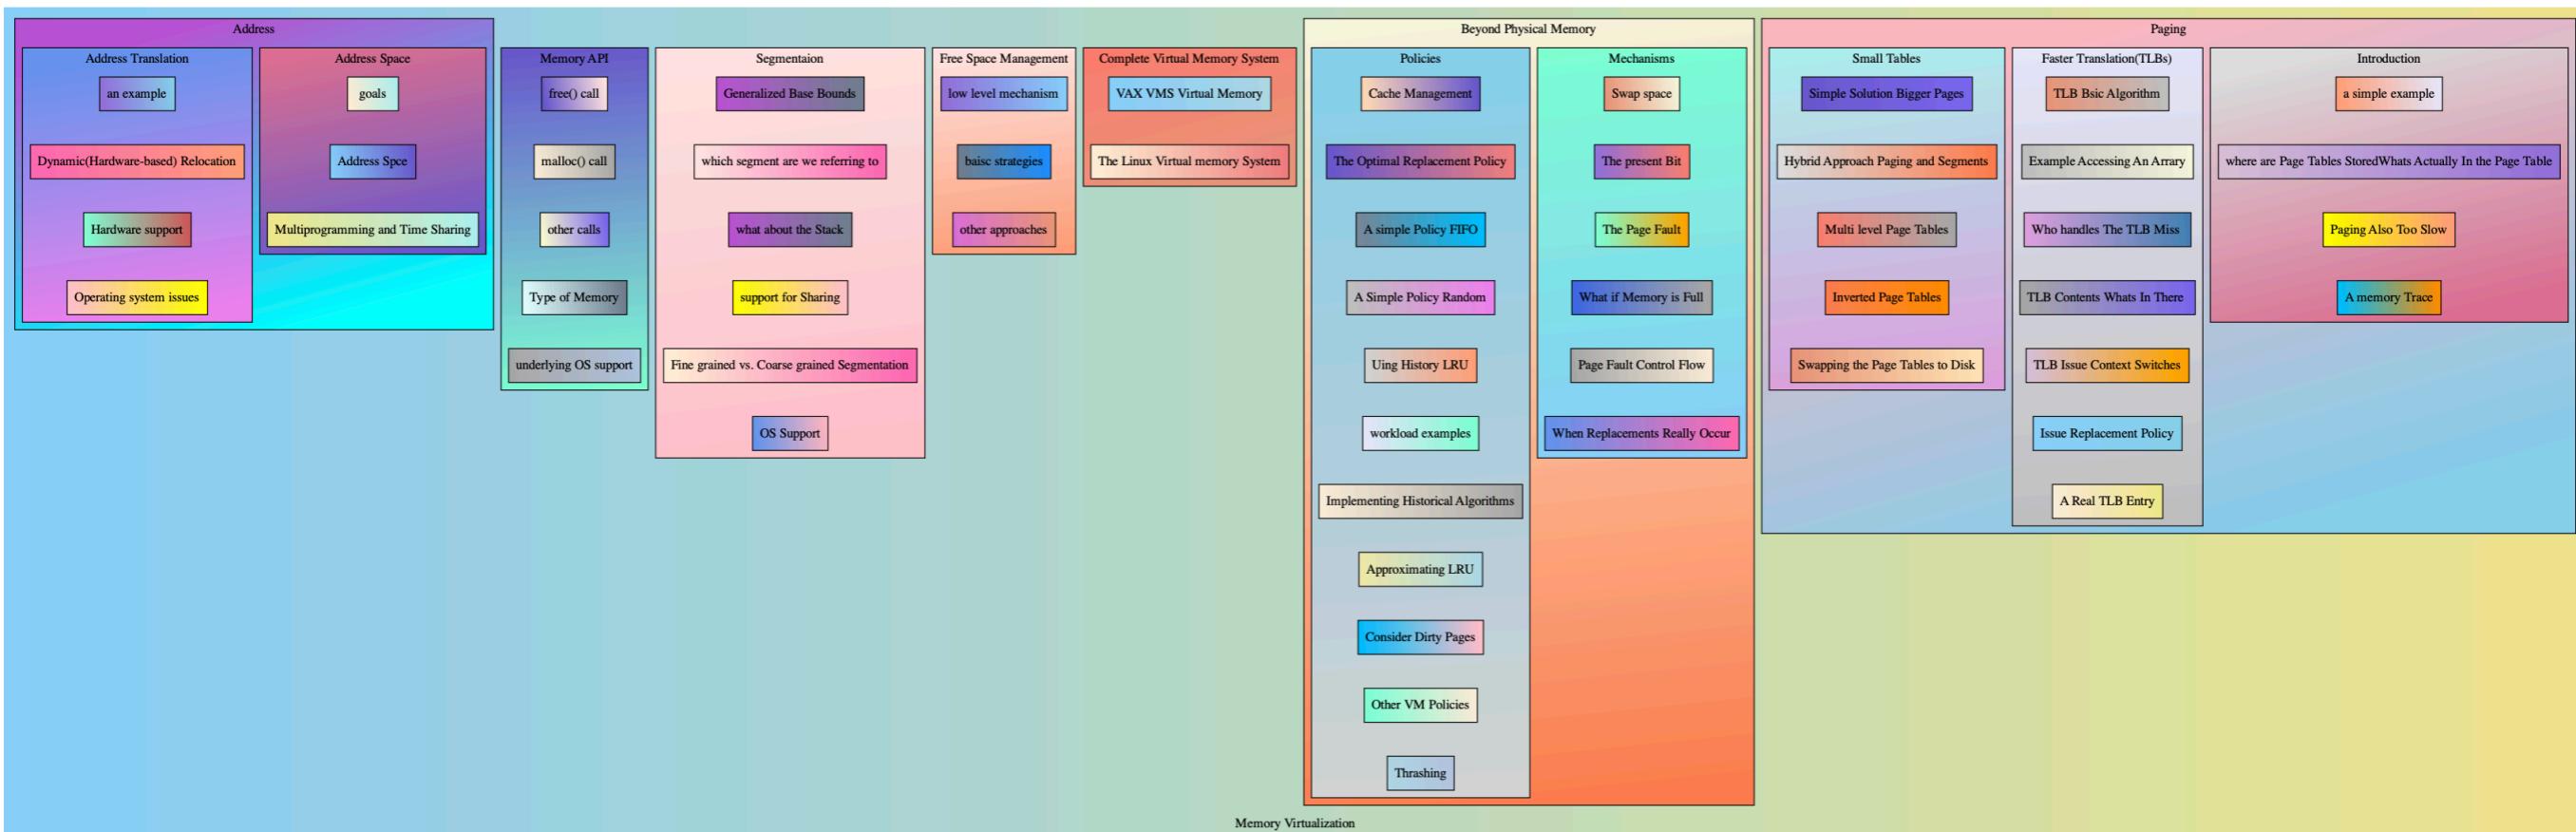
Sun's Network File System(NFS)

The Andrew File System(afs)

Summary Dialogue on Distribution



Memory Virtualization



Concurrency

Introduction

1. Why use Threads?_____
2. How to create a thread?_____
3. Why it get worse, because of shared data?_____
4. The heart of the problem: uncontrolled scheduling
5. One more problem: waiting for another_____
6. The wish for atomicity_____
7. why in OS class?_____

Interlude: Thread API

1. Thread Creation_____
2. Thread Completion_____
3. Locks_____
4. Condition Variables_____
5. Compiling and Running

Locks

1. The basic idea_____
2. Pthread Locks_____
3. Building A Lock_____
4. Evaluating Locks_____
5. Controlling Interrupts_____
6. A Failed Attempt: Just Using Loads Stores_____
7. Building working spin locks with Test-and-set
8. Evaluating Spin Locks_____
9. Compare and swap_____
10. Load-Linked and Store-Conditional_____
11. Fetch-And-Add_____
12. Too much spinning, what now?_____
13. Using Queues: Sleeping instead of Spinning_____
14. Different OS, different support_____
15. Two-Phase Locks_____

Lock-based Concurrent Data Structures

1. Concurrent Counters_____
2. Concurrent Linked Lists
3. Concurrent Queues_____
4. Concurrent Hash Table_____

Condition Variables

1. what is condition variable?_____
2. The routines of condition variable_____
3. The Producer/Consumer(Bounded Buffer) Problem
4. Converging Conditions

Semaphores

1. What is semaphore?_____
2. Binary Semaphores(Locks)_____
3. Semaphores for ordering_____
4. The Producer/Consumer(Bound Buffer) Problem
5. Reader-Writer Locks_____
6. The Dining Philosophers_____
7. How to implement Semaphores?_____

Common Concurrency Problems

1. What types of bugs exist?_____
2. Non-Deadlock bugs_____
3. Deadlock Bugs_____

Event-based Concurrency(Advanced)

1. The basic idea: an event loop_____
2. An Important API: select() / poll()_____
3. using select()_____
4. Why Simpler? No locks needed_____
5. A problem: Blocking System Calls_____
6. A Solution: Asynchronous I/O_____
7. Another Problem: State management_____
8. What is still difficult with events

Persistence

I/O Devices <ul style="list-style-type: none">1. System Architecture2. A Canonical Device3. The Canonical Protocol4. Lowering CPU Overhead with the Interrupts5. More Efficient Data Movement with DMA6. Methods of Device Interaction7. Fitting into the OS: The Device Driver8. Case Study: A simple IDE Disk Driver	Hard Disk Drives <ul style="list-style-type: none">1. The Interface2. Basic Geometry3. A Simple Disk Drive4. I/O Time, Doing the math5. Disk Scheduling	Redundant Arrays of Inexpensive Disks(RAIDs) <ul style="list-style-type: none">1. Interface And RAID Internals2. Fault Model3. How to Evaluate A RAID4. RAID Level 0: Striping5. RAID Level 1: Mirroring6. RAID Level 4: Saving Space with Parity7. RAID Level 5: Rotating Parity8. RAID Comparison: A Summary9. Other Interesting RAID Issues	Interlude: Files and Directories <ul style="list-style-type: none">1. Files and Directories2. The File System Interface3. Creating Files4. Reading and Writing Files5. Reading and Writing, but not Sequentially6. Shared File Table Entries: fork() and dup()7. Writing Immediately with fsync()8. Renaming Files9. Getting Information about Files10. Removing Files11. Making/Reading/Deleting/ Directories12. Hard Links13. Symbolic Links14. Permission Bits and Access Control Lists15. Making and Mounting a file system	File System Implementation <ul style="list-style-type: none">1. The way to think2. Overall Organization3. File Organization: The Inode4. Directory Organization5. Free Space Management6. Access Paths: Reading and Writing7. Caching and Buffering
Locality and The fast file system <ul style="list-style-type: none">1. The Problem: Poor Performance2. FFS: Disk Awareness is the solution3. Organizing Structure: The Cylinder Group4. Policies: How to Allocate Files and Directories5. Measuring File Locality6. The Large-File Exception7. A Few Other things about FFS	Crash Consistency: FSCK and Journaling <ul style="list-style-type: none">1. A Detailed Example2. Solution #1: The File System Checker3. Solution #2: Journaling(or write-Ahead Logging)4. Solution #3: Other Approaches	Log-structured File Systems <ul style="list-style-type: none">1. Writing to Disk Sequentially2. Writing Sequentially And Effectively3. How much to Buffer?4. Problem: Finding Inodes5. Solution Through Indirection: The Inode Map6. Completing The Solution: The Checkpoint Region7. Reading a file from Disk: A recap8. What about directories?9. A new problem: Garbage Collection10. Determining Block Liveness11. A Policy Question: Which Blocks to clean, and when?12. Crash Recovery And the log	Flash-based SSDs <ul style="list-style-type: none">1. Storing a single bit2. From bits to banks/planes3. Basic Flash Operations4. Flash Performance And Reliability5. From Raw Flash to Flash-Based SSDs6. FTL Organization: A bad Approach7. A log-structured FTL8. Garbage Collection9. Mapping Table Size10. Wear Leveling11. SSD Performance and cost	Data Integrity and Protection <ul style="list-style-type: none">1. Disk Failure Modes2. Handling Latent Sector Errors3. Detecting Corruption: The Checksum4. Using Checksums5. A new Problem: Misdirected Writes6. One last Problem: Lost Writes7. Scrubbing8. Overheads of Checksumming
Persistence				